

DIFFERENTIALLY PRIVATE MEDIAN REGRESSION

*E Chen*¹ *Yang Cao*² *Yu Tang*³ *Ying Miao*⁴

¹ Zhejiang Lab ² Hokkaido University ³ Soochow University ⁴ Tsukuba University

ABSTRACT

In signal processing, ensuring privacy protection is of utmost importance due to the handling of sensitive data obtained from various sensors and devices, such as audio recordings, images, and biometric features. In addition, as a reliable approach to address the impact of outliers and extreme values, median regression serves multiple purposes in signal processing, including signal recovery, denoising, parameter estimation, prediction, and interpolation. However, there is no study on enforcing formal privacy for median regression. In this work, we develop differentially private median regression for the first time. Three innovative algorithms, namely DP-SGD, DP-Smooth and DP-ItSq, that provide differential privacy guarantees for median regression, are proposed. Among these algorithms, experimental results consistently demonstrate that DP-Smooth performs remarkably close to non-private methods. Under the premise of protecting privacy, it exhibits high accuracy and delivers the best performance in the experiments, with a relative loss ratio of approximately 99.89%. As for DP-ItSq, it serves as a valuable supplement designed specifically for large sample cases, showcasing efficient computation speed.¹

Index Terms— Differential Privacy, Median Regression, Signal Processing

1. INTRODUCTION

Signal processing applications that involve user-related data, such as face recognition and personalized recommendations, have specifically raised privacy concerns [1]. This has highlighted the critical need to protect users' information from potential breaches and misuse. Therefore, it is crucial to address these privacy concerns and prioritize the safeguarding of user data. Due to the importance of privacy protection in signal processing, the consideration of privacy computation becomes particularly necessary. Differential Privacy (DP) [2] provides a solution to the paradox of obtaining useful information about a population while learning nothing about an individual's specific data. It has gained attention in the field of machine learning [3] and has found real-world applications in preserving privacy [4].

Although median regression is widely recognized as an important method in signal processing [5], the research on privacy protection aspects of this technique remains insufficient. Numerous studies have specifically focused on developing statistical methods that can effectively handle data with low noise. Some notable examples include logistic regression [6], linear regression [7], and kernel ridge regression [8]. However, in situations where there are too many outliers present, it becomes necessary to adopt a robust approach to effectively handle noise. Median regression is a robust statistical method that exhibits resilience to extreme values, outliers, and skewed data distributions [9, 10]. This attribute makes it particularly valuable in signal processing, where such scenarios are prevalent [11]. Moreover, the existing framework for privacy-preserving methods in empirical risk minimization [12] relies on the assumption that the objective function is doubly differentiable and strictly convex. Unfortunately, this assumption does not hold in the case of median regression. Consequently, there is a clear necessity for developing specialized privacy-preserving algorithms that are specifically designed to address the unique requirements of median regression.

In this study, we present three privacy-preserving algorithms specifically designed for median regression: DP-SGD, DP-Smooth and DP-ItSq. To address the non-differentiable nature of the absolute value function at the origin, DP-SGD employs directional derivatives as a substitute for gradients. This algorithm serves as the reference for our research. DP-Smooth, on the other hand, utilizes finite smoothing techniques to handle the non-differentiability issue of the absolute value function [13]. As a complement to DP-Smooth, DP-ItSq converts the median regression problem into a weighted least squares problem [14]. This transformation greatly improves computation speed and achieves good performance when dealing with large a sample size or high privacy budget.

Numerical results show that DP-SGD exhibits significant errors, which do not decrease rapidly with increasing sample size and privacy budget. In contrast, DP-Smooth performs close to non-private algorithms and remains stable across different datasets. As a complement to DP-Smooth, DP-ItSq achieves the fastest computation speed and provides good accuracy, particularly in scenarios involving large a sample size.

¹The full version can be found on website <https://github.com/EChen233/A-full-version-of-Median-regression>.

2. BACKGROUND AND DEFINITIONS

Definition 2.1. $((\epsilon, \delta)$ -DP) A randomized algorithm M is called (ϵ, δ) -indistinguishable if for all $\mathcal{R} \subseteq \text{Range}(M)$ and for all neighboring databases \mathbf{S}, \mathbf{S}' :

$$\mathbb{P}(M(\mathbf{S}) \in \mathcal{R}) \leq e^\epsilon \mathbb{P}(M(\mathbf{S}') \in \mathcal{R}) + \delta, \quad (1)$$

where \mathbf{S} and \mathbf{S}' are neighboring if they differ by exactly one record, and we denote this relationship as $\mathbf{S} \nabla \mathbf{S}'$.

Furthermore, a powerful statistical tool called Gaussian Differential Privacy (GDP) has been proposed [15] to analyze the privacy of algorithms, based on hypothesis testing [16]. It covers the traditional definition of differential privacy and is more general.

Let H_0 represent the underlying dataset as \mathbf{S} , and H_1 represent the underlying dataset as \mathbf{S}' . By denoting U and V as the probabilities of H_0 and H_1 respectively, privacy protection can be defined using the Type I error (α) and Type II error (β) associated with a rejection rule ϕ .

Definition 2.2. For any two probability distributions U and V on the same space Ω , the trade-off function $T(U, V) : [0, 1] \rightarrow [0, 1]$ is defined as

$$T(U, V)(\alpha) = \inf\{\beta_\phi : \alpha_\phi \leq \alpha\}.$$

Definition 2.3 (μ -GDP). A mechanism M is said to satisfy μ -Gaussian Differential Privacy (μ -GDP) if it is G_μ -DP. That is, $T(M(\mathbf{S}), M(\mathbf{S}')) \geq G_\mu$ for all $\mathbf{S} \nabla \mathbf{S}'$. Here $G_\mu = T(N(0, 1), N(\mu, 1))$ for $\mu \geq 0$. An explicit expression for G_μ reads $G_\mu(\alpha) = \Phi(\Phi^{-1}(1 - \alpha) - \mu)$, where $\Phi(\cdot)$ denotes the standard normal CDF.

GDP precisely describes the Gaussian mechanism using a single mean parameter “ μ ” from a unit-variance Gaussian distribution, simplifying the interpretation of privacy guarantees. Indeed, GDP provides a privacy bound that is lossless, meaning that it preserves privacy without any degradation or loss of privacy guarantees. Next, we will introduce a randomized algorithm called the Gaussian mechanism, which is an effective method for privacy preservation.

Definition 2.4 (l_1 sensitivity). The l_1 sensitivity of a function f that outputs a vector in \mathbb{R}^k is:

$$\Delta f = \max_{\mathbf{S} \nabla \mathbf{S}'} \|f(\mathbf{S}) - f(\mathbf{S}')\|_1.$$

Definition 2.5. Given any function f that outputs a vector in \mathbb{R}^k , the Gaussian mechanism is defined as:

$$M(\mathbf{X}, f(\cdot), \mu) = f(\mathbf{S}) + \mathbf{Z},$$

where \mathbf{Z} is drawn from the normal distribution $N(\mathbf{0}, (\frac{\Delta f}{\mu})^2 \mathbf{I}_k)$. The density function of the standard k dimensional normal distribution (centered at 0 with covariance matrix \mathbf{I}_k) is: $\varphi(\mathbf{x}) = (\frac{1}{\sqrt{2\pi}})^k e^{-\frac{\mathbf{x}^T \mathbf{x}}{2}}$.

In this work, we consider median regression and minimize the objective function $L(\beta_0, \beta)$, which combines the sum of absolute residuals $F(\beta_0, \beta)$ with a ridge penalty term:

$$L(\beta_0, \beta) = F(\beta_0, \beta) + \frac{\lambda}{2} \beta^T \beta. \quad (2)$$

The ridge penalty term is determined by the regularization parameter λ and the squared magnitude of the coefficients β .

3. ALGORITHMS

In this section, we provide three privacy preserving algorithms: DP-SGD, DP-Smooth and DP-ItSq, for median regression and calculate their privacy parameters respectively.

3.1. Algorithm 1: DP-SGD

Inspired by [12], we apply a similar approach to minimize the non-differentiable objective function $L(\beta_0, \beta)$, utilizing its directional derivatives. As an illustration, if \mathbf{e}_k represents the coordinate direction of β_k variation, it gives rise to two directional derivatives: $d_{\mathbf{e}_k^+} L(\beta_0, \beta)$ and $d_{\mathbf{e}_k^-} L(\beta_0, \beta)$. We present a theorem regarding privacy, while demonstrating conclusions about utility through simulation results. Following the approach in [17], we utilize a greedy coordinate descent algorithm to update the direction of parameter β_k . This update is based on the equation $Dir^{(k)} = \min\{d_{\mathbf{e}_k^+} L(\beta_0, \beta), d_{\mathbf{e}_k^-} L(\beta_0, \beta)\}$. We terminate the update of β_k if both coordinate directional derivatives are nonnegative.

Algorithm 1 : DP-SGD

Input: Privacy parameters μ , design matrix \mathbf{X} , response vector \mathbf{Y} , regularization parameter λ , learning rate η , privacy parameter within one iteration μ_0 and the number of iterations N_0

Output: $\hat{\omega} \in \mathbb{R}^{d+1}$.

- 1: Initialize the algorithm with a vector $(\hat{\beta}_0(0), \hat{\beta}^T(0))^T$
 - 2: **for** $t = 0, 1, \dots, N_0 - 1$ **do**
 - 3: Subsampling: take a uniformly random subsample $I_t \subseteq [n]$ with batch size m .
 - 4: **for** $i \in I_t$ **do**
 - 5: **Compute gradient:**
 $\mathbf{g}^i(t) = Dir^i(t + 1)$
 - 6: **Clip to norm C :**
 $\tilde{\mathbf{g}}^i(t) = \mathbf{g}^i(t) / \max(C, \|\mathbf{g}^i(t)\|_2)$
 - 7: **end for**
 - 8: $\hat{\beta}(t + 1) = \hat{\beta}(t) - \eta(\frac{1}{m} \sum_{i \in [n]} \tilde{\mathbf{g}}^i(t) + \mathbf{U}(t))$,
 $\mathbf{U}(t) \sim N(\mathbf{0}, \frac{4C^2}{m^2 \mu_0^2} \mathbf{I}_d)$. μ_0 is the value that satisfies
 $\mu = \sqrt{2c} \sqrt{e^{\mu_0^2} \Phi(1.5\mu_0) + 3\Phi(-0.5\mu_0) - 2}$ and $c = m\sqrt{N_0}/n$.
 - 9: $\hat{\beta}_0(t + 1) = \frac{1}{m} \sum_{i=1}^m (Y_i - \mathbf{X}_i \hat{\beta}(t + 1))$.
 - 10: **end for**
 - 11: **return** $\hat{\omega} = (\hat{\beta}_0(N_0), \hat{\beta}(N_0)^T)^T$
-

Theorem 3.1. Given a set of n samples $\mathbf{X}_1, \dots, \mathbf{X}_n$ over \mathbb{R}^d with labels Y_1, \dots, Y_n , where for each i ($1 \leq i \leq n$), $\|\mathbf{X}_i\|_1 \leq 1$ and $|Y_i| \leq B$, the output of DP-SGD preserves μ -GDP.

3.2. Algorithm 2: DP-Smooth

Although DP-SGD can provide privacy protection, due to the nature of the absolute value function, it only includes information about the direction of the gradient during gradient updates, resulting in a decrease in algorithm accuracy. Therefore, we need to carefully design corresponding algorithms, one of which is to use techniques based on finite smoothing [13]. The finite smoothing method is an important tool to solve non-differentiable problem. In addition, the solution of smooth function can estimate the solution of the original function well. This idea is applied in DP-Smooth by an analogous technique. Let γ be a non-negative parameter which indicates the degree of approximation. Define

$$\rho_\gamma(x) = \begin{cases} x^2/(2\gamma), & \text{if } |x| \leq \gamma, \\ |x| - \frac{1}{2}\gamma, & \text{if } |x| > \gamma. \end{cases} \quad (3)$$

Then the nondifferentiable function $L(\beta_0, \beta)$ is approximated by $L_\gamma(\beta_0, \beta) = \frac{1}{n} \sum_{i=1}^n \rho_\gamma(r_i(\beta_0, \beta)) + \frac{\lambda}{2} \beta^T \beta$, where $r_i(\beta_0, \beta)$ is the residual of i -th observation.

Algorithm 2 : DP-Smooth

Input: Privacy parameter μ , design matrix \mathbf{X} , response vector \mathbf{Y} , regularization parameter λ and approximation parameter γ

Output: $\omega^* \in \mathbb{R}^{d+1}$

- 1: Generate a random vector \mathbf{b} from $d+1$ dimensional normal distribution with mean $\mathbf{0}$ and covariance $\frac{16}{\mu^2} \mathbf{I}_{d+1}$
 - 2: **Compute** $(\beta_0^*, \beta^{*T})^T = \underset{\beta_0, \beta}{\operatorname{argmin}} L_\gamma(\beta_0, \beta) + \frac{\mathbf{b}^T \omega}{n} + \frac{\beta_0^2}{\sqrt{n}}$, where $\omega = (\beta_0, \beta^T)^T$ is a $d+1$ dimensional column vector.
 - 3: **return** $\omega^* = (\beta_0^*, \beta^{*T})^T$
-

This algorithm bears a striking resemblance to the smoothing median regression convex program [18], resulting in similar running times as smoothing regression. In fact, ω^* can be obtained by using the interior point method. By following a similar proof as presented in [19], we can demonstrate that DP-Smooth preserves privacy and holds utility.

Theorem 3.2. Given a set of n samples $\mathbf{X}_1, \dots, \mathbf{X}_n$ over \mathbb{R}^d , with labels Y_1, \dots, Y_n , where for each i , $\|\mathbf{X}_i\|_1 \leq 1$ and $|Y_i| \leq B$. If for each i , $r_i(\mu^*, \beta^*) \neq \gamma$, then the output of DP-Smooth would preserve μ -GDP.

Theorem 3.3. Given an l_1 regression problem with regularization parameter λ , let ω_1 be the classifier that minimizes $L_\gamma(\beta_0, \beta) + \frac{\beta_0^2}{\sqrt{n}}$, and ω_2 be the classifier output by DP-smooth

respectively. Then, with probability $1 - \alpha$, $\|\omega_1 - \omega_2\|_1 \leq \frac{4(d+1)t_{1-\alpha}(d+1)}{n\mu \min(\lambda, \frac{2}{\sqrt{n}})}$. Here, $t_{1-\alpha}(d+1)$ represents the $1-\alpha$ quantile of the t -distribution with $d+1$ degrees of freedom.

It should be noted that this is simply a theoretical lower bound on accuracy. In practical simulations, the value of lambda can be set to 0.

3.3. Algorithm 3: DP-ItSq

When the sample size is large, DP-Smooth may suffer from slow computation speed. Therefore, in the case of large samples, we provide a fast algorithm called DP-ItSq. DP-ItSq, which stands for **d**ifferentially **p**rivate **i**terative least **s**quares regression, merges the methodologies of least absolute deviations regression and least squares regression. By doing so, it effectively converts a complex problem involving least absolute deviations into a simple weighted least squares regression [14]. For the $(t+1)$ -th iteration, we set w_i as $\frac{1}{|r(t)_i|+e}$, with $r(t)_i$ being the residual of the i -th sample at the t -th iteration. This leads to the iterative process expressed as:

$$I(t+1) = \frac{1}{n} \sum_{i=1}^n \frac{1}{|r(t)_i|+e} r^2(t+1)_i + \frac{\lambda}{2} \beta^T \beta. \quad (4)$$

If $|r(t)_i - r(t+1)_i| \approx 0, i = 1, 2, \dots, n$, $I(t+1)$ is close to $L(\beta_0, \beta)$. In practice, we set e as a small positive value. According to the relationship between constrained optimization and regularization terms, there exists a mutual correspondence between v and λ , where $\beta^T \beta \leq v$.

Algorithm 3 : DP-ItSq

Input: Privacy parameter μ , design matrix \mathbf{X} , response vector \mathbf{Y} , regularization parameter λ , tolerance parameter τ and the number of iteration N_0

Output: $\hat{\omega} \in \mathbb{R}^{d+1}$

- 1: Partition the dataset into N_0 disjoint subsets, numbered from 0 to $N_0 - 1$.
 - 2: Initialize the algorithm with $\hat{\beta}_0(0)$ and $\hat{\beta}(0)$
 - 3: $(\hat{\beta}_0(1), \hat{\beta}^T(1))^T = \underset{\beta_0, \beta}{\operatorname{argmin}} I(1)$
 - 4: **for** $t = 1, \dots, N_0 - 1$ **do**
 - 5: Estimate the parameters using the t -th subset.
 - 6: **while** $|\hat{\beta}_0(t) - \hat{\beta}_0(t-1)| > \tau$ or $\|\hat{\beta}(t) - \hat{\beta}(t-1)\|_1 > \tau$ **do**
 - 7: $(\hat{\beta}_0(t+1), \hat{\beta}^T(t+1))^T = \underset{\beta_0, \beta}{\operatorname{argmin}} I(t+1)$
 - 8: **end while**
 - 9: **end for**
 - 10: **return** $\hat{\omega} := (\hat{\beta}_0(N_0), \hat{\beta}^T(N_0))^T + \mathbf{U}$, where \mathbf{U} is a $d+1$ dimensional normal random variable with mean $\mathbf{0}$ and covariance $\frac{c^2}{\mu^2} \mathbf{I}_{d+1}$, where $c = \frac{12(\sqrt{dv}+B)}{n \min(\frac{2}{2(\sqrt{dv}+B)+e}, \lambda)e}$.
-

Theorem 3.4. Given a set of n samples $\mathbf{X}_1, \dots, \mathbf{X}_n$ over \mathbb{R}^d , with labels Y_1, \dots, Y_n , where for each i ($1 \leq i \leq n$), $\|\mathbf{X}_i\|_1 \leq 1, |Y_i| \leq B$, the output of DP-ItSq preserves μ -GDP.

Theorem 3.5. Given an l_1 regression problem with a full-rank covariance matrix \mathbf{X} and regularization parameter λ , let ω_1 be the unique minimizer of $L_e(\beta_0, \beta)$ and $\omega_2 = \omega_1 + \mathbf{U}$, where \mathbf{U} is a $d + 1$ dimensional normal random variable with mean $\mathbf{0}$ and covariance $\frac{e^2}{\mu^2} \mathbf{I}_{d+1}$, where c is $\frac{12(\sqrt{dv}+B)}{\min(\frac{2}{2(\sqrt{dv}+B)+e}, \lambda)ne}$. Then, with probability $1 - \alpha$, $\|\omega_1 - \omega_2\|_1 \leq \frac{c}{\mu} \sqrt{d} + 1t_{1-\alpha}(d + 1)$.

4. EXPERIMENTAL RESULTS

Synthetic dataset We constructed the synthetic dataset by sampling $\mathbf{x}_i \in \mathbb{R}^3$, for $i \in [n]$, independently from a multivariate normal distribution with mean $\bar{\mathbf{x}} = (0.2, 0.6, 0.3)$ and covariance $\Sigma = \mathbf{I}_3$. Each \mathbf{x}_i is associated with a corresponding y_i that is generated as $y_i = 0.2 - 3x_{i,1} + 0.5x_{i,2} - x_{i,3} + u_i$, where u_i is sampled from a standard normal distribution $N(0, 0.5)$. In practice, we set the hyperparameters in three algorithms as listed in Table 1. Figure 1 demonstrates that DP-Smooth achieves the highest accuracy across different values of n . The comparison results are not surprising due to the fact that the derivative of DP-SGD only captures sign information, and the accuracy of DP-ItSq improves rapidly with the increase in the sample size.

Table 1. Hyperparameters set of algorithms

Hyperparameters	Value	Description
μ	1	Privacy budget
λ	0.02	Regularized parameter
p	0.02	Sample ratio $p = m/n$ in Alg 1
N_0	100, 957	The number of iterations in Alg 1, 3
η	0.01	Learning rate in Alg 1
μ_0	0.5	Privacy budgets per iteration in Alg 1
γ	0.05	Smoothing parameter in Alg 2
e	0.05	Correcting parameter in Alg 3
v	1	Constraint parameter in Alg 3
B	3	Bound of the response variable
τ	10^{-6}	Tolerance parameter in Alg 3

Carbon nanotubes UCI dataset We obtained the data from a dataset investigating atomic coordinates in carbon nanotubes [20]. Although this dataset does not include sensitive information, we utilize it to assess the performance of differential privacy algorithms on diverse real datasets. Once we removed all points that do not lie within the interval $[0, 1] \times [0, 1]$, the resulting dataset comprises a total of 10,721 datapoints. All parameters are set the same as Table 1, except for μ . Figure 2 demonstrates that as the privacy budget increases, DP-ItSq converges rapidly to a vicinity of the optimal solution, while DP-SGD shows no significant change. DP-Smooth exhibits the best fitting performance at the cost of higher computational complexity.




Fig. 1. Utility for various sample size.

Fig. 2. Utility for various privacy budgets.




Fig. 3. Comparing fitting accuracy of ECG signal with DP-Smooth under various privacy budgets

Electrocardiogram signal ECG signals, which represent the electrical activity of the heart, are frequently affected by substantial noise and outliers. In this study, we utilize data from the “ECG-ID” database, a part of the PhysioNet biological signal bank [21], which consists of 310 ECG recordings obtained from 90 individuals. Considering the findings from previous analysis, we employ DP-Smooth for fitting purposes. Given that the ECG signal is a periodic sequence, we employ a 6-th order Fourier series expansion [22] with respect to the independent variable “Time”. That is, $Y = a_0 + \sum_{i=1}^6 (a_i \cos(i\omega t) + b_i \sin(i\omega t))$, where $\omega = \frac{2\pi}{T}$, with period $T = \frac{5}{6}$ s. Figure 3 indicates that $\mu = 2$ is already a sufficiently good fit and the smaller the degree of privacy protection, the better the fitting effect. This can be observed from the consistency of peak values between the original sequence and the fitted sequence.

5. CONCLUSION

In this study, three differential privacy algorithms, namely DP-SGD, DP-Smooth, and DP-ItSq, are introduced as solutions for median regression. Among these algorithms, DP-Smooth demonstrated high accuracy in the experiments, with a relative loss ratio of approximately 99.89%. Additionally, DP-ItSq served as a valuable supplement specifically designed for large sample cases, showcasing efficient computation. In the future, our objective is to investigate privacy protection algorithms for various signal processing tasks, including the processing of image data.

6. REFERENCES

- [1] Reginald L Lagendijk, Zekeriya Erkin, and Mauro Barni, “Encrypted signal processing for privacy protection: Conveying the utility of homomorphic encryption and multiparty computation,” *IEEE Signal Processing Magazine*, vol. 30, no. 1, pp. 82–105, 2012.
- [2] Cynthia Dwork, Aaron Roth, et al., “The algorithmic foundations of differential privacy,” *Foundations and Trends® in Theoretical Computer Science*, vol. 9, no. 3–4, pp. 211–407, 2014.
- [3] Martin Abadi, Andy Chu, Ian Goodfellow, H Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang, “Deep learning with differential privacy,” in *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, 2016, pp. 308–318.
- [4] Úlfar Erlingsson, Vasyl Pihur, and Aleksandra Korolova, “Rappor: Randomized aggregatable privacy-preserving ordinal response,” in *Proceedings of the 2014 ACM SIGSAC conference on computer and communications security*, 2014, pp. 1054–1067.
- [5] Jose L Paredes and Gonzalo R Arce, “Compressive sensing signal reconstruction by weighted median regression estimates,” *IEEE Transactions on Signal Processing*, vol. 59, no. 6, pp. 2585–2601, 2011.
- [6] Ø Birkenes, Tomoko Matsui, Kunio Tanabe, Sabato Marco Siniscalchi, Tor André Myrvoll, and Magne Hallstein Johnsen, “Penalized logistic regression with hmm log-likelihood regressors for speech recognition,” *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 18, no. 6, pp. 1440–1454, 2010.
- [7] T Tony Cai, Yichen Wang, and Linjun Zhang, “The cost of privacy: Optimal rates of convergence for parameter estimation with differential privacy,” *The Annals of Statistics*, vol. 49, no. 5, pp. 2825–2850, 2021.
- [8] Jie Chen, Lingfei Wu, Kartik Audhkhasi, Brian Kingsbury, and Bhuvana Ramabadrari, “Efficient one-vs-one kernel ridge regression for speech recognition,” in *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2016, pp. 2454–2458.
- [9] Friedrich-Wilhelm Scholz, “Weighted median regression estimates,” *The Annals of Statistics*, vol. 6, no. 3, pp. 603–609, 1978.
- [10] Keming Yu, Zudi Lu, and Julian Stander, “Quantile regression: applications and current research areas,” *Journal of the Royal Statistical Society Series D: The Statistician*, vol. 52, no. 3, pp. 331–350, 2003.
- [11] Alan T Welford, “Signal, noise, performance, and age,” *Human Factors*, vol. 23, no. 1, pp. 97–109, 1981.
- [12] Di Wang, Minwei Ye, and Jinhui Xu, “Differentially private empirical risk minimization revisited: Faster and more general,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [13] Gianni Di Pillo, Luigi Grippo, and Stefano Lucidi, “A smooth method for the finite minimax problem,” *Mathematical Programming*, vol. 60, no. 1-3, pp. 187–214, 1993.
- [14] EJ Schlossmacher, “An iterative technique for absolute deviations curve fitting,” *Journal of the American Statistical Association*, vol. 68, no. 344, pp. 857–859, 1973.
- [15] Jinshuo Dong, Aaron Roth, and Weijie J Su, “Gaussian differential privacy,” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, vol. 84, no. 1, pp. 3–37, 2022.
- [16] Peter Kairouz, Sewoong Oh, and Pramod Viswanath, “The composition theorem for differential privacy,” in *International conference on machine learning*. PMLR, 2015, pp. 1376–1385.
- [17] Jerome Friedman, Trevor Hastie, Holger Höfling, and Robert Tibshirani, “Pathwise coordinate optimization,” *The Annals of Applied Statistics*, vol. 1, no. 2, pp. 302 – 332, 2007.
- [18] Marcelo Fernandes, Emmanuel Guerre, and Eduardo Horta, “Smoothing quantile regressions,” *Journal of Business & Economic Statistics*, vol. 39, no. 1, pp. 338–357, 2021.
- [19] Kamalika Chaudhuri and Claire Monteleoni, “Privacy-preserving logistic regression,” *Advances in neural information processing systems*, vol. 21, 2008.
- [20] Mehmet Acı and Mutlu Avcı, “Artificial neural network approach for atomic coordinate prediction of carbon nanotubes,” *Applied Physics A*, vol. 122, pp. 1–14, 2016.
- [21] Tatiana S Lugovaya, “Biometric human identification based on electrocardiogram,” *Master’s thesis, Faculty of Computing Technologies and Informatics, Electrotechnical University ‘LETI’, Saint-Petersburg, Russian Federation*, 2005.
- [22] Gerald B Folland, *Fourier analysis and its applications*, vol. 4, American Mathematical Soc., 2009.
- [23] David R Hunter and Kenneth Lange, “A tutorial on mm algorithms,” *The American Statistician*, vol. 58, no. 1, pp. 30–37, 2004.

7. APPENDIX

We present the following necessary lemmas, and the detailed proofs can be found in the article [15].

Lemma 7.1. *The Gaussian mechanism is μ -GDP.*

Lemma 7.2 (μ -GDP composition). *The k -fold composition (applying the same algorithm to a dataset k times) of μ_0 -GDP mechanisms is $\sqrt{k}\mu_0$ -GDP.*

Lemma 7.3. *If M is f -DP on X^m , then the subsampled mechanism $M(\text{Sample}_m)$ is $C_p(f)$ -DP on X^n , where the sampling ratio $p = m/n$. Especially, if $m\sqrt{N_0}/n = c$ and $f = G_{\mu_0}$, then $C_p(G_{\mu_0})^{\otimes N_0}$ is asymptotically μ -GDP with*

$$\mu = \sqrt{2c} \sqrt{e^{\mu_0^2} \Phi(1.5\mu_0) + 3\Phi(-0.5\mu_0) - 2}.$$

7.1. Details of DP-SGD

The computation of directional gradients is crucial in this context. Specifically, the directional derivative of β_k can be computed as follows.

$$\begin{aligned} d_{\mathbf{e}_k^+} L(\beta_0, \boldsymbol{\beta}) &= \lim_{\tau \rightarrow 0^+} \frac{L(\beta_0, \boldsymbol{\beta} + \tau \mathbf{e}_k) - L(\beta_0, \boldsymbol{\beta})}{\tau} \\ &= d_{\mathbf{e}_k^+} F(\beta_0, \boldsymbol{\beta}) + \lambda \beta_k, \end{aligned}$$

and

$$\begin{aligned} d_{\mathbf{e}_k^-} L(\beta_0, \boldsymbol{\beta}) &= \lim_{\tau \rightarrow 0^-} \frac{L(\beta_0, \boldsymbol{\beta} + \tau \mathbf{e}_k) - L(\beta_0, \boldsymbol{\beta})}{\tau} \\ &= d_{\mathbf{e}_k^-} F(\beta_0, \boldsymbol{\beta}) + \lambda \beta_k. \end{aligned}$$

In l_1 regression, the coordinate direction derivatives are

$$d_{\mathbf{e}_k^+} F(\beta_0, \boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n \begin{cases} -x_{ik}, & r_i(\beta_0, \boldsymbol{\beta}) < 0, \\ x_{ik}, & r_i(\beta_0, \boldsymbol{\beta}) > 0, \\ |x_{ik}|, & r_i(\beta_0, \boldsymbol{\beta}) = 0, \end{cases} \quad (5)$$

and

$$d_{\mathbf{e}_k^-} F(\beta_0, \boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n \begin{cases} -x_{ik}, & r_i(\beta_0, \boldsymbol{\beta}) < 0, \\ x_{ik}, & r_i(\beta_0, \boldsymbol{\beta}) > 0, \\ -|x_{ik}|, & r_i(\beta_0, \boldsymbol{\beta}) = 0. \end{cases} \quad (6)$$

Proof of Theorem 3.1

Proof. For $(x, y) \in (\mathbf{X}(t), \mathbf{Y}(t))$, it suffices to prove the privacy guarantee for the t -th iteration of the algorithm and use μ -GDP composition to obtain full privacy bound. At the t -th iteration, the algorithm first updates the non-sparse estimate of $\boldsymbol{\beta}$:

$$\hat{\boldsymbol{\beta}}(t+1) = \hat{\boldsymbol{\beta}}(t) - \eta \left(\frac{1}{m} \sum_{i \in [m]} \tilde{\mathbf{g}}^i(t) + \mathbf{U}(t) \right),$$

where $\mathbf{U}(t) \sim N\left(\mathbf{0}, \frac{4C^2}{m^2\mu_0^2} \mathbf{I}_d\right)$. For convenience, assume the first $m-1$ records are the same, then the sensitivity of $\frac{1}{m} \sum_{i \in [m]} \tilde{\mathbf{g}}^i(t) \leq \frac{2C}{m}$.

Combined with Lemma 7.1, $\hat{\boldsymbol{\beta}}(t+1)$ satisfies μ_0 -GDP. In addition, since $\hat{\mu} = \frac{1}{n_0} \sum_{i=1}^{n_0} (Y_i - \mathbf{X}_i \hat{\boldsymbol{\beta}})$, it is differentially private by post-processing. By using Lemma 7.3, DP-SGD is asymptotically μ -GDP, where

$$\mu = \sqrt{2c} \sqrt{e^{\mu_0^2} \Phi(1.5\mu_0) + 3\Phi(-0.5\mu_0) - 2}.$$

□

7.2. Details of DP-Smooth

Since the absolute value function is not differentiable at the cuspidal point, a smooth method for minimizing function (2) is considered. Denote $F_\gamma(\beta_0, \boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n \rho_\gamma(r_i(\beta_0, \boldsymbol{\beta}))$, and $L_\gamma(\beta_0, \boldsymbol{\beta}) = F_\gamma(\beta_0, \boldsymbol{\beta}) + \frac{\lambda}{2} \boldsymbol{\beta}^T \boldsymbol{\beta}$. The sign vector

$$\mathbf{S}_\gamma(\beta_0, \boldsymbol{\beta}) = (s_1(\beta_0, \boldsymbol{\beta}), \dots, s_n(\beta_0, \boldsymbol{\beta}))^T$$

is given by

$$s_i(\beta_0, \boldsymbol{\beta}) = \begin{cases} -1, & \text{if } r_i(\beta_0, \boldsymbol{\beta}) < -\gamma, \\ 0, & \text{if } -\gamma \leq r_i(\beta_0, \boldsymbol{\beta}) \leq \gamma, \\ 1, & \text{if } r_i(\beta_0, \boldsymbol{\beta}) > \gamma. \end{cases} \quad (7)$$

Let $w_i(\beta_0, \boldsymbol{\beta}) = 1 - s_i^2(\beta_0, \boldsymbol{\beta})$, then

$$\begin{aligned} \rho_\gamma(r_i(\beta_0, \boldsymbol{\beta})) &= \frac{1}{2\gamma} w_i(\beta_0, \boldsymbol{\beta}) r_i^2(\beta_0, \boldsymbol{\beta}) \\ &\quad + s_i(\beta_0, \boldsymbol{\beta}) \left[r_i(\beta_0, \boldsymbol{\beta}) - \frac{1}{2} \gamma s_i(\beta_0, \boldsymbol{\beta}) \right]. \end{aligned}$$

Denote $\mathbf{W}_\gamma(\beta_0, \boldsymbol{\beta})$ as the diagonal $n \times n$ matrix whose diagonal elements are $w_i(\beta_0, \boldsymbol{\beta})$. So $\mathbf{W}_\gamma(\beta_0, \boldsymbol{\beta})$ has value 1 in those diagonal elements related to small residuals and 0 elsewhere. For $\beta_0 \in \mathbb{R}$ and $\boldsymbol{\beta} \in \mathbb{R}^d$, the derivation of $F_\gamma(\beta_0, \boldsymbol{\beta})$ is

$$\frac{\partial F_\gamma(\beta_0, \boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = \frac{1}{n} \mathbf{X}^T \left[\frac{1}{\gamma} \mathbf{W}_\gamma(\beta_0, \boldsymbol{\beta}) r(\beta_0, \boldsymbol{\beta}) + \mathbf{S}_\gamma(\beta_0, \boldsymbol{\beta}) \right],$$

and

$$\frac{\partial F_\gamma(\beta_0, \boldsymbol{\beta})}{\partial \beta_0} = \frac{1}{n} \mathbf{1}^T \left[\frac{1}{\gamma} \mathbf{W}_\gamma(\beta_0, \boldsymbol{\beta}) r(\beta_0, \boldsymbol{\beta}) + \mathbf{S}_\gamma(\beta_0, \boldsymbol{\beta}) \right].$$

Notice that $L_\gamma(\beta_0, \boldsymbol{\beta})$ is strictly convex and differentiable. By intuition, a minimizer of $L(\beta_0, \boldsymbol{\beta})$ is close to a minimizer of $L_\gamma(\beta_0, \boldsymbol{\beta})$ when γ is close to zero. Furthermore, the l_1 solution can be detected when $\gamma > 0$ is small enough [13]. That is, it is not necessary to let γ converge to zero in order to find a minimizer of $L_\gamma(\beta_0, \boldsymbol{\beta})$. This observation is essential for the efficiency and the numerical stability of DP-Smooth.

Proof of Theorem 3.2

Proof. Let \mathbf{a}_1 and \mathbf{a}_2 be two row vectors over \mathbb{R}^d with l_1 norm at most 1 and $y_1, y_2 \in [-B, B]$. Consider the two inputs D_1 and D_2 where D_2 is obtained from D_1 by replacing one record (\mathbf{a}_1, y_1) into (\mathbf{a}_2, y_2) . For convenience, assume the first $n - 1$ records are the same. For any output ω^* by Algorithm 1, there is a unique value of \mathbf{b} that maps the input to the output. This uniqueness holds, because both the regularization function and the loss function are differentiable everywhere, and the objective function is strictly convex. Denote $\tilde{\mathbf{a}}_1$ as $(1, \mathbf{a}_1)$ and $\tilde{\mathbf{a}}_2$ as $(1, \mathbf{a}_2)$, w_n and w'_n are corresponding coefficients. In addition, let \mathbf{X} and \mathbf{X}' represent samples about independent variables of D_1 and D_2 , and \mathbf{Y} and \mathbf{Y}' represent their corresponding response vectors, respectively. Denote $\tilde{\mathbf{X}} = (1, \mathbf{X})$ and $\tilde{\mathbf{X}}' = (1, \mathbf{X}')$, where $\mathbf{1}$ is an n dimensional column vector whose elements are 1. Let the values of $d + 1$ dimensional vector \mathbf{b} for D_1 and D_2 respectively, be \mathbf{b}_1 and \mathbf{b}_2 . Since ω^* is the value that minimizes both the optimization problems, the derivative of both optimization functions at ω^* is 0. This implies that for every \mathbf{b}_1 in the first case, there exists a \mathbf{b}_2 in the second case such that:

$$\begin{aligned} & \mathbf{b}_1 + \tilde{\mathbf{X}}^T (1/\gamma \mathbf{W}_\gamma(\mu^*, \beta^*) (\tilde{\mathbf{X}}\omega^* - \mathbf{Y}) + \mathbf{S}_\gamma(\mu^*, \beta^*)) \\ &= \mathbf{b}_2 + \tilde{\mathbf{X}}'^T (1/\gamma \mathbf{W}'_\gamma(\mu^*, \beta^*) (\tilde{\mathbf{X}}'\omega^* - \mathbf{Y}') + \mathbf{S}'_\gamma(\mu^*, \beta^*)), \end{aligned}$$

which indicates that:

$$\begin{aligned} & \mathbf{b}_1 + \tilde{\mathbf{a}}_1^T (1/\gamma w_n(\mu^*, \beta^*) (\mu^* + \mathbf{a}_1\beta^* - y_1) + s_n(\mu^*, \beta^*)) \\ &= \mathbf{b}_2 + \tilde{\mathbf{a}}_2^T (1/\gamma w'_n(\mu^*, \beta^*) (\mu^* + \mathbf{a}_2\beta^* - y_2) + s'_n(\mu^*, \beta^*)). \end{aligned}$$

According to the definitions of $\mathbf{W}_\gamma(\mu^*, \beta^*)$ and $\mathbf{S}_\gamma(\mu^*, \beta^*)$, it is clear that

$$-1 \leq 1/\gamma w_n(\mu^*, \beta^*) * (\mu^* + \mathbf{a}_1\beta^* - y_1) + s_n(\mu^*, \beta^*) \leq 1$$

and

$$-1 \leq 1/\gamma w'_n(\mu^*, \beta^*) * (\mu^* + \mathbf{a}_2\beta^* - y_2) + s'_n(\mu^*, \beta^*) \leq 1.$$

Since $\|\tilde{\mathbf{a}}_1\|_1 \leq 2$ and $\|\tilde{\mathbf{a}}_2\|_1 \leq 2$, we have $\|\mathbf{b}_1 - \mathbf{b}_2\|_1 \leq 4$ and $-4 \leq \|\mathbf{b}_1\|_1 - \|\mathbf{b}_2\|_1 \leq 4$, which implies that DP-smooth approximately satisfies μ -GDP. \square

In theory, for each i , the probability of ω^* satisfies $|r_i(\mu^*, \beta^*)| \neq \gamma$ is 1.

Lemma 7.4. [19] Let $G(\omega)$ and $g(\omega)$ be two convex functions, which are continuous and differentiable at all points. If $\omega_1 = \operatorname{argmin}_\omega G(\omega)$ and $\omega_2 = \operatorname{argmin}_\omega G(\omega) + g(\omega)$, then $\|\omega_1 - \omega_2\|_1 \leq \frac{g_1}{G_2}$. Here, $g_1 = \max_\omega \|\nabla g(\omega)\|_1$ and $G_2 = \min_v \min_\omega \mathbf{v}^T \nabla^2 G(\omega) \mathbf{v}$, for any unit vector \mathbf{v} .

The main idea of the proof is to examine the gradient and the Hessian of the functions G and g around ω_1 and ω_2 . For fixed \mathbf{b} , let $\omega = (\beta_0, \beta)$, $G(\omega) = L_\gamma(\beta_0, \beta)$ and $g(\omega) = \frac{\beta_0^2}{\sqrt{n}} + \frac{\beta_0^2}{\sqrt{n}}$, we can obtain accuracy of DP-Smooth. Notice that \mathbf{b} is a random variable, it is necessary to provide its property of convergence.

Lemma 7.5. If \mathbf{b} is a random variable drawn from $d + 1$ dimensional normal distribution with mean $\mathbf{0}$ and covariance \mathbf{I}_{d+1} , then $\|\mathbf{b}\|_2$ obeys t distribution with parameter $d + 1$.

Proof. Since $\mathbf{b} = (b_1, \dots, b_{d+1})' \sim N(\mathbf{0}, \mathbf{I}_{d+1})$, then for $1 \leq i < j \leq d + 1$, b_i and b_j are independent and identically distributed as $N(0, 1)$. Hence, $\|\mathbf{b}\|_2^2$ is the sum of $d + 1$ independent standard normal variables, which follows a $\chi^2(d + 1)$ distribution. By the definition of t distribution, the result is obtained. \square

Proof of Theorem 3.3

Proof. According to Lemma 7.4, we take $G(\omega) = L_\gamma(\beta_0, \beta) + \frac{\beta_0^2}{\sqrt{n}}$ and $g(\omega) = \frac{\beta^T \omega}{n}$. Because $F_\gamma(\beta_0, \beta)$ is a convex function, if we define the second derivative of $F_\gamma(\beta_0, \beta)$ to be 0 at nondifferentiable points, then the Hessian matrix of $F_\gamma(\beta_0, \beta)$ is positive semi-definite. Notice that

$$\nabla^2 \left(\frac{\beta_0^2}{\sqrt{n}} \right) = \begin{bmatrix} \frac{2}{\sqrt{n}} & \mathbf{0}_{1 \times d} \\ \mathbf{0}_{d \times 1} & \mathbf{0}_{d \times d} \end{bmatrix}_{(d+1) \times (d+1)},$$

and

$$\nabla^2 \left(\frac{\lambda}{2} \beta^T \beta \right) = \begin{bmatrix} 0 & 0 \\ 0 & \lambda \mathbf{I}_{d \times d} \end{bmatrix}_{(d+1) \times (d+1)}.$$

Hence, for any unit vector \mathbf{v} , we have $G_2 \geq \min(\lambda, \frac{2}{\sqrt{n}})$ and $g_1 = \frac{\|\mathbf{b}\|_1}{n}$, $\|\omega_1 - \omega_2\|_1 \leq \frac{\|\mathbf{b}\|_1}{n \min(\lambda, \frac{2}{\sqrt{n}})} \leq \frac{(d+1)\|\mathbf{b}\|_2}{n \min(\lambda, \frac{2}{\sqrt{n}})}$. Combined with Lemma 7.5, the theorem is obtained. \square

7.3. Details of DP-ItSq

For $e > 0$, define a perturbation of $L(\beta_0, \beta)$ as

$$L_e(\beta_0, \beta) = \sum_{i=1}^n |r_i(\beta_0, \beta)| - \frac{e}{2} \log(e + |r_i(\beta_0, \beta)|) + \frac{\lambda}{2} \beta^T \beta.$$

[23] proved that iterative least squares algorithm without adding noise is a special case of Majorization-Minimization (MM) algorithm for the objective function $L_e(\beta_0, \beta)$ and obtained the following convergence results.

Proposition 7.1. For linear median regression with a full-rank covariate matrix \mathbf{X} , the iterative least squares algorithm without adding noise converges to the unique minimizer of $L_e(\beta_0, \beta)$.

Proposition 7.2. If $(\hat{\beta}_{0,e}, \hat{\beta}_e)$ minimizes $L_e(\beta_0, \beta)$, then any limit point of $(\hat{\beta}_{0,e}, \hat{\beta}_e)$ as e tends to 0 minimizes $L(\beta_0, \beta)$. If $L(\beta_0, \beta)$ has a unique minimizer $(\tilde{\beta}_0, \tilde{\beta})$, then $\lim_{e \rightarrow 0} (\hat{\beta}_{0,e}, \hat{\beta}_e) = (\tilde{\beta}_0, \tilde{\beta})$.

Proposition 7.1 and 7.2 show that $(\hat{\beta}_0(N_0), \hat{\beta}^T(N_0))^T$ can estimate true value well. Notice that DP-ItSq is simply adding normal noise to true parameters, hence, the accuracy of output in DP-ItSq is given directly.

Proof of Theorem 3.4

Proof. Denote $\hat{\omega} = (\hat{\beta}_0(N_0), \hat{\beta}^T(N_0))^T$ and the l_1 sensitivity of $\hat{\omega}$ as $s(\hat{\omega})$. Let \mathbf{a}_1 and \mathbf{a}_2 be two row vectors over \mathbb{R}^d with l_1 norm at most 1 and $y_1, y_2 \in [-B, B]$. Consider the two inputs D_1 and D_2 where D_2 is obtained from D_1 by changing one record (\mathbf{a}_1, y_1) into (\mathbf{a}_2, y_2) . For convenience, assume the first $n-1$ records are the same and denote $\tilde{\mathbf{a}}_1$ as $(1, \mathbf{a}_1)$ and $\tilde{\mathbf{a}}_2$ as $(1, \mathbf{a}_2)$. According to Lemma 7.4, let $G(\hat{\omega}) = I(N_0)$ and $g(\hat{\omega}) = \frac{1}{n}w_n(\tilde{\mathbf{a}}_2\hat{\omega} - y_2)^2 - \frac{1}{n}w_n(\tilde{\mathbf{a}}_1\hat{\omega} - y_1)^2$. Similar to the proof in Theorem 3.2, we can achieve that

$$\|\nabla g(\hat{\omega})\|_1 = \left\| \frac{\partial g(\omega)}{\partial \beta_0} \right\|_{\omega=\hat{\omega}} + \left\| \frac{\partial g(\omega)}{\partial \beta} \right\|_{\omega=\hat{\omega}},$$

where

$$\frac{\partial g(\omega)}{\partial \beta_0} \Big|_{\omega=\hat{\omega}} = \frac{2w_n}{n} [(\mathbf{a}_2 - \mathbf{a}_1)\hat{\beta}(N_0) - (y_2 - y_1)],$$

$$\begin{aligned} \frac{\partial g(\omega)}{\partial \beta} \Big|_{\omega=\hat{\omega}} &= \frac{2w_n}{n} [(\hat{\beta}_0(N_0) + \mathbf{a}_2\hat{\beta}(N_0) - y_2)\mathbf{a}_2^T \\ &\quad - (\hat{\beta}_0(N_0) + \mathbf{a}_1\hat{\beta}(N_0) - y_1)\mathbf{a}_1^T]. \end{aligned}$$

According to the triangle inequality, we can achieve that

$$\begin{aligned} \left\| \frac{\partial g(\omega)}{\partial \beta} \right\|_{\omega=\hat{\omega}} &\leq \frac{2}{n} |w_n| (|\hat{\beta}_0(N_0)| + |\mathbf{a}_2\hat{\beta}(N_0)| + |y_2| \\ &\quad + |\hat{\beta}_0(N_0)| + |\mathbf{a}_1\hat{\beta}(N_0)| + |y_1|). \end{aligned}$$

Notice that $(\hat{\beta}_0(N_0), \hat{\beta}^T(N_0))^T = \text{argmin}_{\beta_0, \beta} I(N_0)$, then $\frac{\partial I(N_0)}{\partial \beta_0} = 0$ at $\beta_0 = \hat{\beta}_0(N_0)$, that is,

$$\begin{aligned} \sum_{i=1}^n w_i (\hat{\beta}_0(N_0) + \mathbf{X}_i \hat{\beta}(N_0) - Y_i) &= 0 \\ \iff \hat{\beta}_0(N_0) &= \frac{-\sum_{i=1}^n w_i (\mathbf{X}_i \hat{\beta}(N_0) - Y_i)}{\sum_{i=1}^n w_i}. \end{aligned}$$

Since $0 < w_i \leq 1/e$ ($i = 1, \dots, n$) and $\|\hat{\beta}(N_0)\|_1 \leq \sqrt{d}\|\hat{\beta}(N_0)\|_2 \leq \sqrt{dv}$, then we have $|\hat{\beta}_0(N_0)| \leq \sqrt{dv} + B$ and we obtain the following inequality:

$$\left\| \frac{\partial g(\omega)}{\partial \beta} \right\|_{\omega=\hat{\omega}} \leq \frac{8(\sqrt{dv} + B)}{ne}.$$

Similarly, it can be verified that

$$\left\| \frac{\partial g(\omega)}{\partial \beta_0} \right\|_{\omega=\hat{\omega}} \leq \frac{4(\sqrt{dv} + B)}{ne}.$$

Therefore,

$$g_1 = \max_{\hat{\omega}} \|\nabla g(\hat{\omega})\|_1 \leq \frac{12(\sqrt{dv} + B)}{ne}.$$

Notice that above inequalities are still true in t -th ($t \geq 2$) iteration and hence $\frac{1}{2(\sqrt{dv} + B) + e} \leq w_i \leq \frac{1}{e}$. In addition, denote

$F_e(\omega) = \frac{1}{n} \sum_{i=1}^n w_i r_i^2(\beta_0, \beta)$. It can be checked that $F_e(\omega)$ is convex and

$$\frac{\partial F_e^2(\omega)}{\partial \beta_0^2} = \frac{2}{n} \sum_{i=1}^n w_i \geq \frac{2}{2(\sqrt{dv} + B) + e},$$

$$\frac{\partial^2 (\frac{\lambda}{2} \beta^T \beta)}{\partial \beta} = \lambda \mathbf{I},$$

where \mathbf{I} is an identity matrix with size $d \times d$, then $G_2 \geq \min(\frac{2}{2(\sqrt{dv} + B) + e}, \lambda)$ and $s(\hat{\omega}) \leq \frac{12(\sqrt{dv} + B)}{n \min(\frac{2}{2(\sqrt{dv} + B) + e}, \lambda)e}$.

According to Lemma 7.1, the result is obtained. \square

Proof of Theorem 3.5

Proof. By using Lemma 7.5, $\frac{\mu}{c} \|\mathbf{U}\|_2 \sim t(d+1)$. Then with probability $1 - \alpha$, $\|\mathbf{U}\|_1 \leq \sqrt{d+1} \|\mathbf{U}\|_2 \leq \sqrt{d+1} t_{1-\alpha}(d+1)$, then the theorem is obtained. \square